

Data Structures and Algorithms

Lesson 6: Keep Trees Balanced



AVL trees, red-black trees, TreeSet, TreeMap

Outline

1. **Deletion**
2. Balancing Trees
 - Rotations
3. AVL Trees
4. Red-Black Trees
 - TreeSet
 - TreeMap

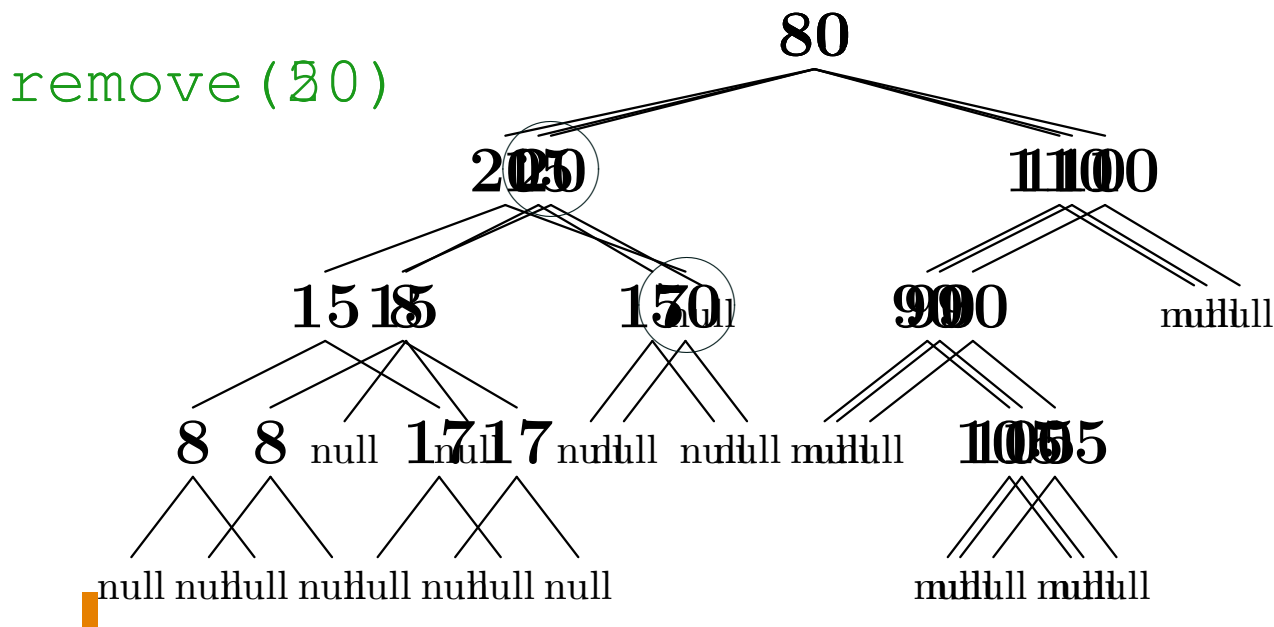


Recap

- Binary search trees are commonly used to store data because we need to only look down one branch to find any element
- We saw how to implement many methods of the binary search tree
 - `contains`
 - `add`
 - `successor` (in outline)
- One method we missed was `remove`

Deletion

- Suppose we want to delete an element from a binary search tree ■
- It is relatively easy if the element is held in a leaf node (e.g. 50) ■
- It is not so hard if the node holding the element has one child (e.g. 20) ■




Code for remove

```

if (e.left==null && e.right==null) {
    // leaf node
    if (e == e.parent.left)
        e.parent.left = null;
    else
        e.parent.right = null;
} else if (e.right==null) {
    // no right child
    if (e == e.parent.left)
        e.parent.left = e.left;
    else
        e.parent.right = e.left;
    e.left.parent = e.parent;
} else if (e.left==null) {
    // no left child
    if (e == e.parent.left)
        e.parent.left = e.right;
    else
        e.parent.right = e.right;
    e.right.parent = e.parent;
}

```




Removing Element with Two Children

■ If an element has two children then ■

- replace that element by its successor ■
- and then remove the successor using the above procedure ■

remove(80)




Outline

1. Deletion
2. **Balancing Trees**
 - Rotations
3. AVL Trees
4. Red-Black Trees
 - TreeSet
 - TreeMap




Why Balance Trees

- The number of comparisons to access an element depends on the depth of the node
- The average depth of a node depends on the shape of the tree



full tree



sparse tree


- The shape of the tree depends on the order the elements are added

Tree Depth

- In the best case (a full tree), if the number of elements in the tree is $n = 2^l - 1$ then the maximum depth of a node is $l = \log_2(n + 1)$, which is $\Theta(\log(n))$
- In the worst case (when the tree is effectively a linked list), the maximum depth of a node is n , which is $\Theta(n)$
- It turns out for random sequences the average depth of a node is $O(\log(n))$
- Unfortunately, the worst case happens when the elements are added *in order* (not a rare event)

Rotations


- To avoid unbalanced trees we would like to modify the shape
- This is possible as the shape of the tree is not uniquely defined (e.g. we could make any node the root)
- We can change the shape of a tree using **rotations**
- E.g. left rotation



Types of Rotations

■ We can get by with 4 types of rotations


- Left rotation (as above)
- Right rotation (symmetric to above)



- Left-right double rotation
- Right-left double rotation


Coding Rotations

```
void rotateLeft(Node<T> e){  
    // r = right node  
    Node<T> r = e.right;  
  
    // link e with its new right node  
    e.right = r.left;  
    if (r.left != null)  
        r.left.parent = e;  
  
    // link r to its new parent  
    r.parent = e.parent;  
    if (e.parent == null)  
        root = r;  
    else if (e.parent.left == e)  
        e.parent.left = r;  
    else  
        e.parent.right = r;  
  
    // fix links between r and e  
    r.left = e;  
    e.parent = r;  
}
```




When Single Rotations Work

- Single rotations balance the tree when the unbalanced subtree is on the outside



Double Rotations

- If the unbalanced subtree is on the inside we need a double rotation



```
leftRotation(C.left);  
rightRotation(C);
```

Outline

1. Deletion
2. Balancing Trees
 - Rotations
3. **AVL Trees**
4. Red-Black Trees
 - TreeSet
 - TreeMap



Balancing Trees


- There are different strategies for using rotations for balancing trees
- The three most popular are
 - AVL trees
 - Red-black trees
 - Splay trees
- They differ in the criteria they use for doing rotations

AVL Trees

- AVL trees were invented in 1962 by two Russian mathematicians Adelson-Velski and Landis
- An AVL tree is a binary search tree in which
 1. the heights of the left and right subtree differ by at most 1
 2. the left and right subtrees are AVL trees
- This guarantees that the worst case AVL tree has logarithmic depth

Minimum Number of Nodes

- We want to see how full an AVL tree has to be, at the minimum
- Let $m(h)$ be the minimum number of nodes in an AVL tree of height h
- This has to be made up of a root and two subtrees: one of height $h - 1$; and, in the worst case, one of height $h - 2$



- Thus, the least number of nodes in an AVL tree of height h is

$$m(h) = m(h - 1) + m(h - 2) + 1$$

- with $m(1) = 1$, $m(2) = 2$

Proof of Exponential Number of Nodes

- We have $m(h) = m(h - 1) + m(h - 2) + 1$,
with $m(1) = 1, m(2) = 2$
- This gives us a sequence $1, 2, 4, 7, 12, \dots$
- Compare this with Fibonacci $f(h) = f(h - 1) + f(h - 2)$,
with $f(1) = f(2) = 1$
- This gives us a sequence $1, 1, 2, 3, 5, 8, 13, \dots$
- It looks like $m(h) = f(h + 2) - 1$
 - this can be proved by induction
- This shows that $m(h)$ grows exponentially with h

Proof of Logarithmic Depth

■ $m(h) = m(h-1) + m(h-2) + 1$ with $m(1) = 1, m(2) = 2$

■ We can prove by induction $m(h) \geq (3/2)^{h-1}$ ■

$$m(1) = 1 \geq (3/2)^0 = 1, \quad m(2) = 2 \geq (3/2)^1 = 3/2 \quad \checkmark \quad \blacksquare$$

$$m(h) = m(h-1) + m(h-2) + 1 \geq \left(\frac{3}{2}\right)^{h-2} + \left(\frac{3}{2}\right)^{h-3} + 1 = \left(\frac{3}{2}\right)^{h-3} \left(\frac{3}{2} + 1 + \left(\frac{3}{2}\right)^{3-h}\right)$$

$$\geq \left(\frac{3}{2}\right)^{h-3} \frac{5}{2} = \left(\frac{3}{2}\right)^{h-3} \frac{10}{4} \geq \left(\frac{3}{2}\right)^{h-3} \frac{9}{4} = \left(\frac{3}{2}\right)^{h-1} \quad \checkmark \quad \blacksquare$$

■ The number of elements, n , we can store in an AVL tree of height h is $n \geq m(h)$, and thus $n \geq (3/2)^{h-1}$ ■

■ Taking logs: $\log(n) \geq (h-1) \log(3/2)$ ■

■ So $h \leq \frac{\log(n)}{\log(3/2)} + 1$ ■

■ That is, h is $O(\log(n))$ ■

Implementing AVL Trees

- To implement an AVL tree, we include additional information at each node indicating the balance of the subtrees:


$$\text{balanceFactor} = (\text{height of left subtree}) - (\text{height of right subtree})$$

- So for an AVL tree:

$$\text{balanceFactor} = \begin{cases} -1 & \text{right subtree deeper than left subtree} \\ 0 & \text{left and right subtrees equal} \\ +1 & \text{left subtree deeper than right subtree} \end{cases}$$


Implementing AVL Trees

- update `balanceFactor` after each `add/remove` operation
- rebalance tree using rotations if needed



Balancing AVL Trees

- When adding an element to an AVL tree
 - Find the location where it is to be inserted, and insert
 - Iterate up through the parents re-adjusting the `balanceFactor`
 - **If** the balance factor exceeds ± 1 then re-balance the tree and stop – why can we stop?
 - **else if** the balance factor goes to zero **then** stop – why can we stop?



AVL Deletions

- AVL deletions are similar to AVL insertions
 - need to update balance factors/rebalance tree
- One difference is that after performing a rotation the tree may still not satisfy the AVL criteria so higher levels need to be examined
- In the worst case $\Theta(\log(n))$ rotations may be necessary
- This may be relatively slow – but in many applications deletions are rare

AVL Tree Performance

- Insertion, deletion and search in AVL trees are, at worst, $\Theta(\log(n))$
 - height of an AVL tree is $\Theta(\log(n))$
 - so searching is at worst $\Theta(\log(n))$
 - insertion without balancing is $\Theta(\log(n))$, balancing takes an additional $\Theta(\log(n))$ steps in the worst case
 - deletion without balancing is $\Theta(\log(n))$ at worst (need to find the node first), balancing takes an additional $\Theta(\log(n))$ steps in the worst case ■
- The height of an average AVL tree is $\approx 1.44 \log_2(n)$
- The height of an average binary search tree is $\approx 2.1 \log_2(n)$ ■
- Despite being more compact, insertion is slightly slower in AVL trees than in BSTs without balancing■
- Search is, of course, quicker!■

Outline

1. Deletion
2. Balancing Trees
 - Rotations
3. AVL Trees
4. **Red-Black Trees**
 - TreeSet
 - TreeMap




Red-Black Trees

- Red-black trees are another strategy for balancing trees
- Nodes are either **red** or **black**
- Two rules ensure that **no path from the root to a leaf, or to a node with one child, is more than twice as long as another:**


Red Rule: the children of a red node must be black

Black Rule: the number of black nodes must be the same in all paths from the root to nodes with no children or with one child



Restructuring

- When inserting a new element we first find its position
- If it is the root we colour it black otherwise we colour it red
- If its parent is red we must either relabel or restructure the tree
 - relabel if the “uncle” exists and is also red
 - rotate if the “uncle” does not exist or it is black



Performance of Red-Black Trees

- Red-black trees are slightly more complicated to code than AVL trees
- Red-black trees tend to be slightly less compact than AVL trees
- However, insertion and deletion run slightly quicker
- Both Java Collection classes and C++ STL use red-black trees

TreeSet

- The Java Collection class has a `TreeSet` class implemented using a red-black tree
- It also has a `HashSet` class (which we cover later)
- The `TreeSet` class iterates over the elements in order (unlike the `HashSet` class)
- These are the classes to use if you want a collection of objects with no repetitions

Maps

- One major abstract data type (ADT) we have not yet seen an implementation for is maps
- The Java `Map` interface contains element pairs `Map<K, V>`
 - The first element of type `K` is the `key`
 - The second element of type `V` is the `value`
- The Java class `TreeMap<K, V>` implements this interface, using Red-Black trees
- Maps work as content addressable arrays

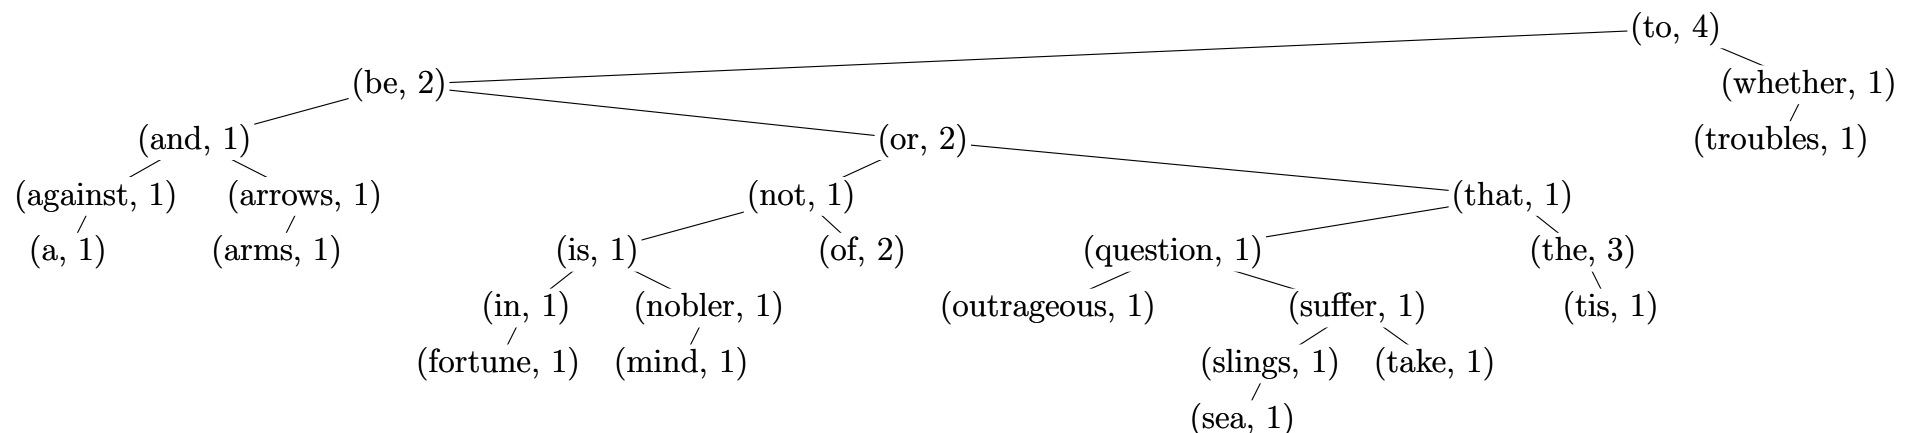
```
Map<String,Integer> students = new TreeMap<String,Integer>();  
student.put("John_Smith", 89);  
student.put("Terry_Jones", 98);  
System.out.println(students.get("John_Smith"));
```


Implementing a TreeMap

- We can implement Map using a TreeSet by making each node hold a `Entry<K, V>` object

```
public class Entry<K, V> implements Comparable
{
    private K key;
    private V value;
}
```

- We can count occurrences of words using the key for words and the value to count



Lessons

- Binary search trees are very efficient (order $\log(n)$ insertion, deletion and search) provided they are balanced
- Balanced trees are achieved by performing rotations
- There are different strategies for deciding when to rotate including
 - AVL trees
 - Red-black trees
- Binary search trees are used for implementing **sets** and **maps**