

B-Tree supplementary slides

ECS, University of Southampton

B-Trees


In the lecture on tries and B-Trees, the insertion process for a B-Tree is illustrated on an example where $M = 5$ and $L = 5$.

There are a number of variants of B-trees. The definition on lecture slide 7 requires all data items to be stored in the leaves (so the internal nodes cannot store data items).


This variation is sometimes known as a B+ tree, although (annoyingly) this term can mean a number of things. In particular, it is commonly understood that a B+ tree allows efficient in-order traversal of the data in the leaves by linking the leaves (this is *not* done in the example from the lecture).

It is perhaps easier to follow the logic behind the insertion procedure for this variant of B-tree in a smaller example.

B-Trees




B-Trees



N.B. In this variant of a B-tree, the data items can only be stored in the *leaves* (so in the root node 41 and 66 are not data items, but they are in the leaves).


B-Trees

insert(30)



B-Trees


insert(30)



41 and 66 in the root act as separators and keys are inserted according to their membership in intervals defined by these separators. In this case $30 < 41$ and there is space to store the item in the first child (leaf) node.


B-Trees

insert(60)



B-Trees


insert(60)



60 lies in the interval $[41, 66)$, and so is stored in the second leaf (again, space permits).


B-Trees

insert(35)



B-Trees

insert(35)



When we insert 35, we run out of space in the first leaf node, which cannot store the data items 11, 23, 30, 35. We split the leaf into two parts 11, 23 and 30, 35; if we are “right-biased”, we pick 30 as our median and insert 30 as a separator into the parent node if we can. If not, we follow the same splitting process with the parent node.